



Chapter 1: Introduction

252-0061-00 V Systems Programming and Computer Architecture

This course covers in *depth*...

- How to write **fast** and **correct** code
- How to write good ***systems*** code
- What makes programs go fast (and slow)
- Programming in C
 - *Still* the systems programming language of choice
- Programming in Assembly Language
 - What the machine understands
- Programs as more than mathematical objects
 - E.g. how does Facebook work?
 - How programs **interact** with the hardware

Who are we?



Prof. Timothy Roscoe

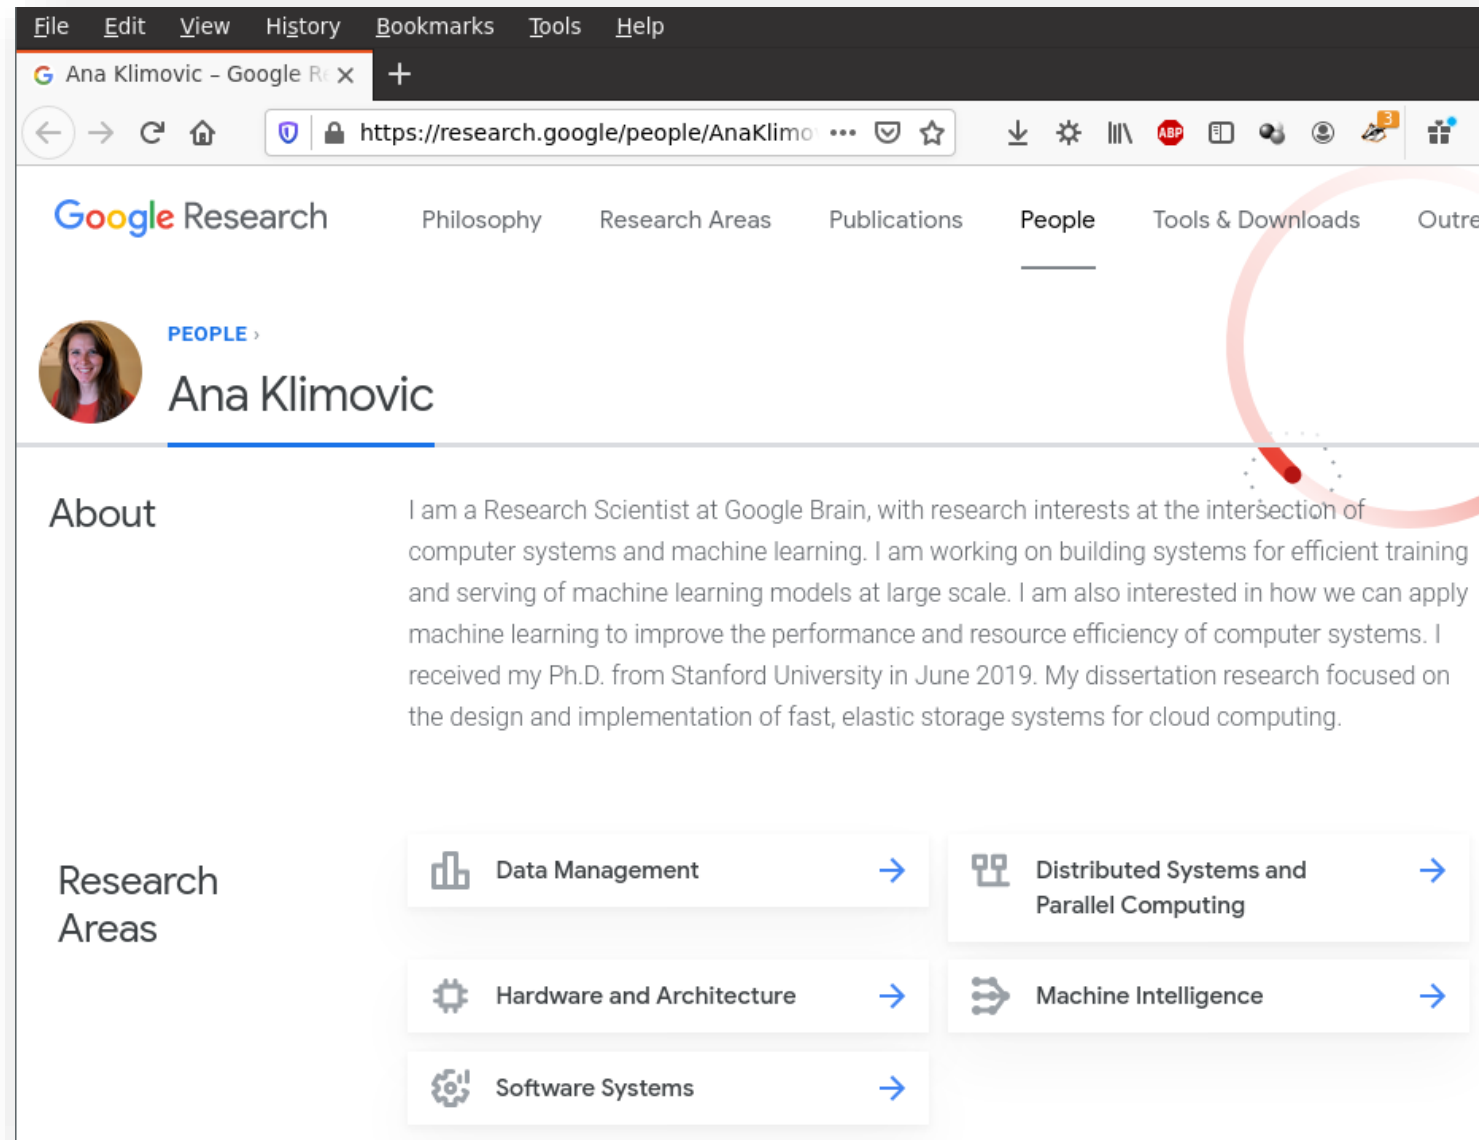


Prof. Ana Klimovic

Full Disclosure



Full Disclosure




The screenshot shows a web browser window with the Google Research profile of Ana Klimovic. The browser's address bar displays the URL `https://research.google/people/AnaKlimovic`. The page header includes navigation links: Google Research, Philosophy, Research Areas, Publications, People (which is underlined), Tools & Downloads, and Outreach. The profile section features a circular profile picture of Ana Klimovic, a blue 'PEOPLE' tag, and her name 'Ana Klimovic'. Below this is the 'About' section, which is circled in red. The 'About' text reads: 'I am a Research Scientist at Google Brain, with research interests at the intersection of computer systems and machine learning. I am working on building systems for efficient training and serving of machine learning models at large scale. I am also interested in how we can apply machine learning to improve the performance and resource efficiency of computer systems. I received my Ph.D. from Stanford University in June 2019. My dissertation research focused on the design and implementation of fast, elastic storage systems for cloud computing.' At the bottom, the 'Research Areas' section lists five categories, each with an icon and a right-pointing arrow: Data Management, Distributed Systems and Parallel Computing, Hardware and Architecture, Machine Intelligence, and Software Systems.

File Edit View History Bookmarks Tools Help

Ana Klimovic - Google Research

Google Research Philosophy Research Areas Publications People Tools & Downloads Outreach

 PEOPLE Ana Klimovic

About

I am a Research Scientist at Google Brain, with research interests at the intersection of computer systems and machine learning. I am working on building systems for efficient training and serving of machine learning models at large scale. I am also interested in how we can apply machine learning to improve the performance and resource efficiency of computer systems. I received my Ph.D. from Stanford University in June 2019. My dissertation research focused on the design and implementation of fast, elastic storage systems for cloud computing.

Research Areas

- Data Management →
- Distributed Systems and Parallel Computing →
- Hardware and Architecture →
- Machine Intelligence →
- Software Systems →

Acknowledgements

- Lots of material from the famous **CS 15-213** at Carnegie Mellon University
 - Basis for the **book**
- Some C programming slides adapted from **CSE333** at University of Washington
 - Many thanks to (ex-)Prof. Steve Gribble
- New material:
 - Considerable evolution...
 - Multicore, devices, etc.
 - Mostly our fault 😊

1.1: Logistics

Systems Programming and Computer Architecture

Lectures

*The slides are **not**
intended to be
understood without
the lectures...*

- Physical:
 - 10:00-12:00 Tuesdays and Wednesdays
 - HG E 7 (Tue) & NO C 60 (Wed)
- Recordings will appear after a few days
 - <https://video.ethz.ch/lectures/d-infk/2023/autumn.html>

Moodle

<https://moodle-app2.let.ethz.ch/course/view.php?id=18098>

- The first place to look!
- **Links** posted here
- All **lecture materials** will be posted on Moodle.
- Ask **questions** in the **forum**
 - TAs and Profs will monitor the forum!
- We will **not** answer questions on Discord.

The screenshot shows a web browser displaying the Moodle course page for '252-0061-00L Systems Programming and Computer Architecture HS2023'. The page has a dark header with the ETH Zürich logo and navigation links like 'Dashboard', 'Kurssuche', 'Support', and 'Neuigkeiten'. A sidebar on the left contains a menu with categories like 'About the course', 'Communication', 'Course hours', 'Logistics', 'Additional material', and 'Lectures'. The main content area on the right is titled '252-0061-00L Systems Programming and Computer Architecture HS2023' and includes tabs for 'Kurs', 'Einstellungen', 'Teilnehmer/innen', and 'Mehr'. Under the 'About the course' section, it describes the course content (introduction to systems programming, C and assembly, floating point arithmetic, etc.) and lists three objectives for students. At the bottom, it mentions the course is based on 'Computer Systems: A Programmer's Perspective' (3rd Edition) by R. Bryant and D. O'Hallaron.

Tutorial sessions

- Very important!
- Logistics:
 - Wednesday, 12:00-14:00 or 14:00-16:00
 - See myStudies for rooms and **streams**
- Content:
 - Tools and skills for **lab exercises**
 - Knowledge needed for exams, but not in the lectures!
- There **will** be a session this Wednesday (tomorrow)

Language

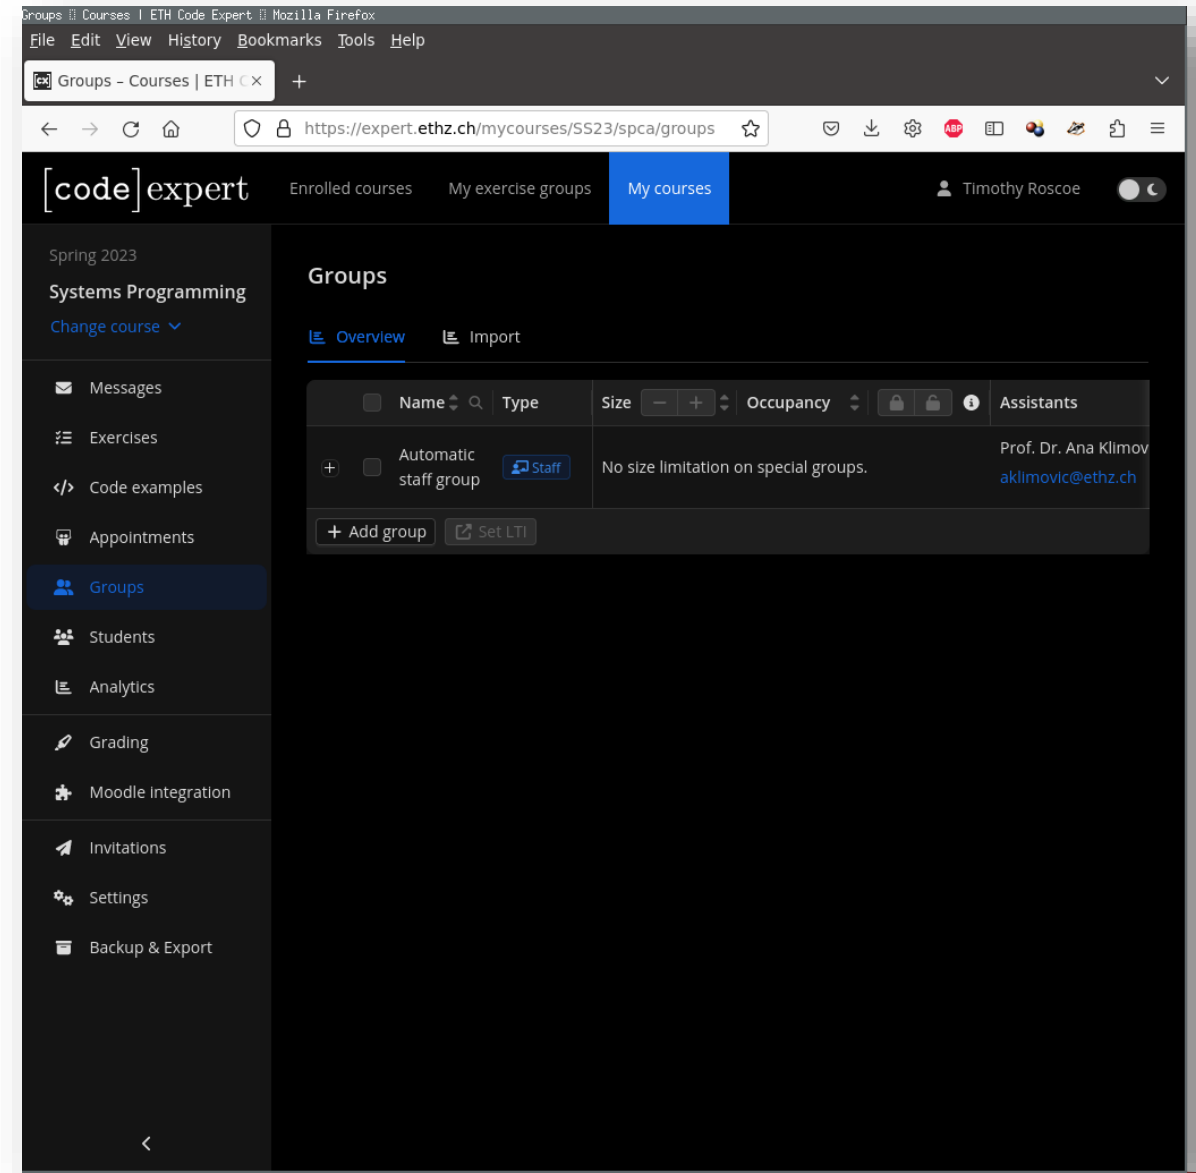
- We'll teach in English (and C...)
 - If we speak too fast, or say something unclear, raise your hand!
 - Please ask questions!
- Assistants speak German, English, Italian, French, ...

Asking questions

1. Ask **during the lectures**
2. Ask on the **Moodle forum** outside the lectures
3. Ask your **friends**
4. Check the **web**
5. Ask your teaching **assistant**
6. Ask ***another*** teaching assistant
7. Email us (troscoe@inf.ethz.ch or aklimovic@ethz.ch)

What's new?

- This year we will use **CodeExpert** for the exam.
 - There will be a **midterm exam**
 - It will **not** contribute to your grade
 - Some exercises, labs and demos will move to CodeExpert
- Course **topics** will **not** change significantly
 - **Mode of answering** questions will
 - **More C experience** will be needed



What's new?


- You'll be seeing a lot more code
 - We'll make all source code available to play with
- You'll need to be comfortable with the command line
 - We'll spend a lot of time at the prompt.
 - Unix shell commands
 - `ls`, `more`, `make`, `gcc`, `gdb`, `grep`, `rm`, `mv`, ...
- Why?
 - Closer to reality: the topic of this course
 - Makes in clearer what is really going on

```
Shell
cixous: ..ing-systems-programming> ls
assignments/      extra_docs/      README.md
Dates.xlsx        lectures/         source_code/
exercise_sessions/ 'Material update.pdf' 'Teaching schedule.pdf'
cixous: ..ing-systems-programming> cd source_code/chapter01/
./source_code/chapter01/
cixous: ..g/source_code/chapter01> ls
addint/  copyij.c  copyji.c  timing.c
cixous: ..g/source_code/chapter01> 
```

Programming environment

- This course targets:
 - CodeExpert (of course)
 - Linux Ubuntu 22.04 LTS on 64-bit x86 PC hardware
- Various options for this exist:
 - **Native** install, or log into lab machines or Optimus
 - **Windows** Subsystem for Linux (from the Store)
 - Virtual machine installation (e.g. **VirtualBox**)
 - *Use another environment and wing it...*
- *We'll help, but there's a limit to how weird we're prepared to get.*

Questions?



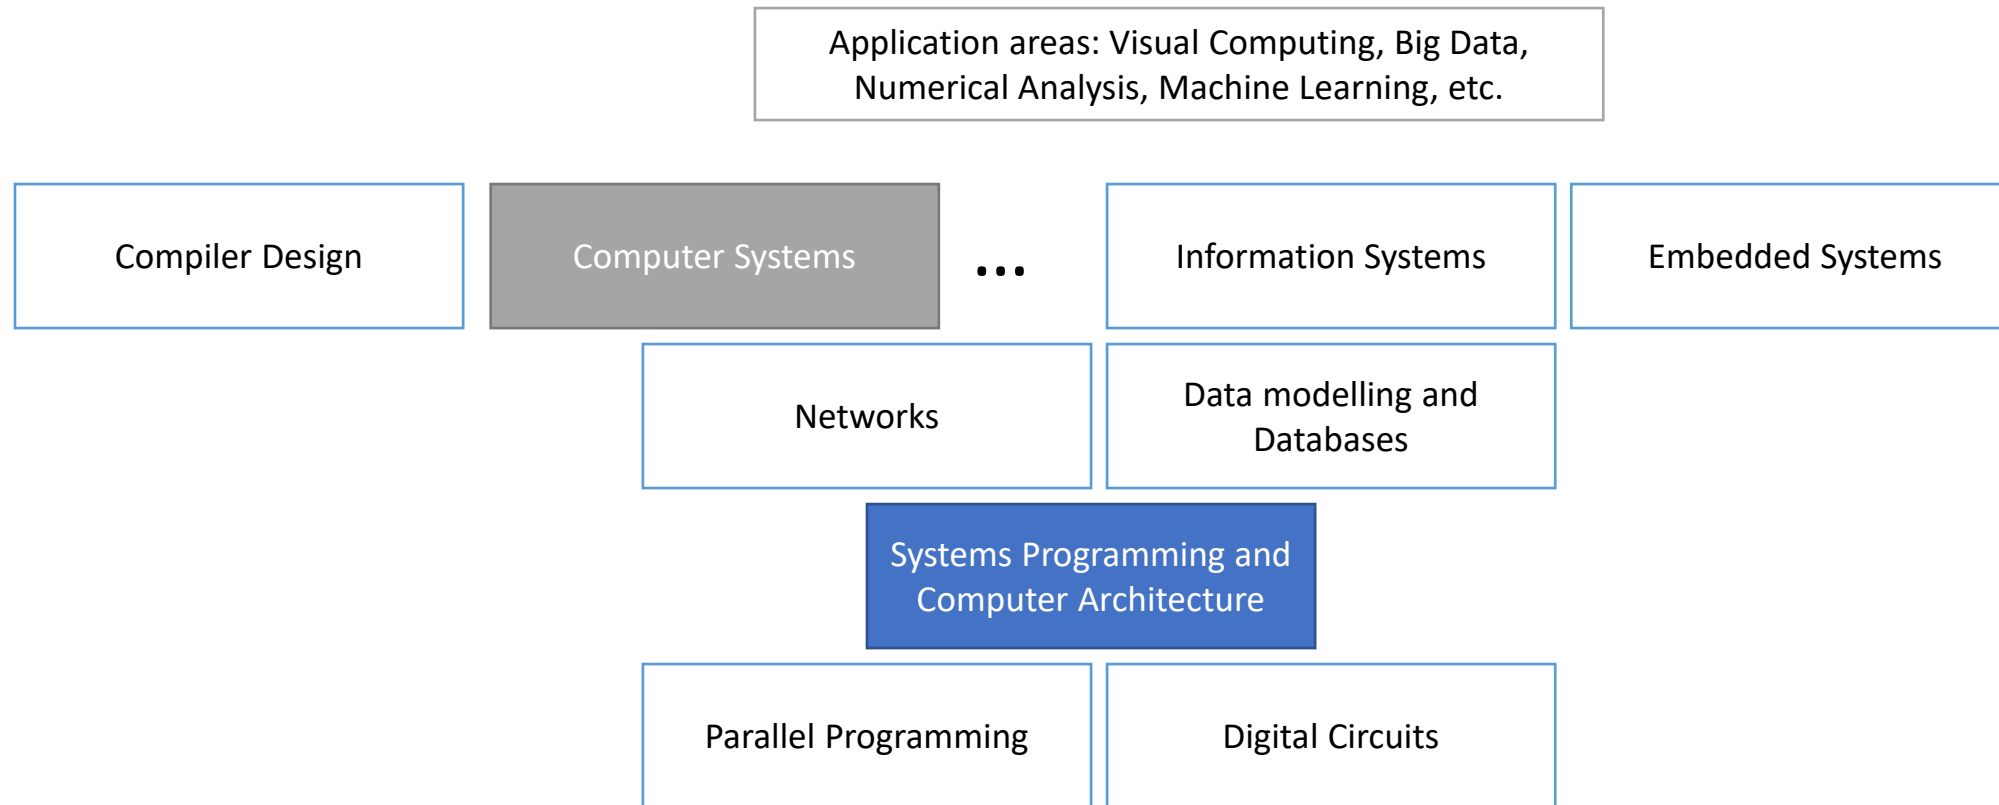
1.2: What is Systems Programming?

Systems Programming and Computer Architecture

“Systems” as a field

- Encompasses:
 - Operating systems
 - Database systems
 - Networking protocols and routing
 - Compiler design and implementation
 - Distributed systems
 - Cloud computing & online services
 - Big Data and machine learning frameworks
- On and above the *hardware/software* boundary

You are here:



“Systems” as a field

„In designing an operating system one needs both theoretical insight and horse sense. Without the former, one designs an ad hoc mess; without the latter one designs an elephant in best Carrara marble (white, perfect, and immobile).“

Roger Needham and David Hartley,
ACM Symposium on Operating Systems Principles,
1968

Motivation

- Most CS courses emphasize **abstraction**
 - Abstract data types (objects, contracts, etc.)
 - Program as mathematical object with well-defined behavior
 - Asymptotic analysis (worst-case, complexity)
- These abstractions have **limitations**
 - Often don't survive contact with reality
 - Especially in the presence of bugs
 - Need to understand details of underlying implementations

Summary: Course Goals

- Become more effective programmers
 - Find and eliminate **bugs** efficiently
 - Understand and tune for program **performance**
- Prepare for later **systems** classes at ETHZ
 - Compilers, Operating Systems, Networks, Computer Architecture, Embedded Systems

Questions?

1.3: Motivation – Some inconvenient truths about computers

Systems Programming and Computer Architecture

Inconvenient truth:
Computers don't really deal
with numbers.

Computers don't deal with integers

- Maths:

$$S = \{i_0 \dots i_{k-1}, i \in \mathbb{Z}\}$$

$$T = \sum_{j=0}^{k-1} i_j$$

- Reality:

```
#include <stdio.h>
#include <stdlib.h>

#define BUFFER_LENGTH 80

int main(int argc, char *argv[])
{
    char buffer[BUFFER_LENGTH];
    int total = 0;

    while( fgets( buffer, BUFFER_LENGTH, stdin ) ) {
        total += atoi(buffer);
    }
    printf("Total is %d\n", total);
    return 0;
}
```

This is not a bug – it is correct behavior!

Computers don't deal with reals either

- Maths:

$$\forall x, y, z \in \mathbb{R},$$

$$(x + y) + z = x + (y + z)$$

- Reality:

```
#include <stdio.h>

int main(int argc, char *argv[])
{
    float x = 1e20;
    float y = -1e20;
    float z = 3.14;

    printf("( x + y ) + z = %f\n", (x + y) + z);
    printf("x + ( y + z ) = %f\n", x + (y + z));
    return 0;
}
```

This is not a bug – it is correct behavior!

Computer arithmetic

- Does not generate random values
 - Arithmetic operations have important mathematical properties
- Cannot assume all “usual” mathematical properties
 - Due to finiteness of representations
 - Integer operations satisfy “ring” properties
 - Commutativity, associativity, distributivity
 - Floating point operations satisfy “ordering” properties
 - Monotonicity, values of signs
- Observation
 - Need to understand which abstractions apply in which contexts
 - Important issues for compiler writers and serious application programmers

Inconvenient truth:

All the best programmers know
assembly.

You've got to know assembly

- Chances are, you'll never write a program in assembly
 - Compilers are much better & more patient than you are
- But: understanding assembly is **key** to machine-level execution model
 - Behavior of programs in presence of **bugs**
 - High-level language model breaks down
 - Tuning program **performance**
 - Understand optimizations done/not done by the compiler
 - Understanding sources of program inefficiency
 - Implementing **system software**
 - Compiler has machine code as target
 - Operating systems must manage process state
 - Creating / fighting **malware**
 - x86 assembly is the language of choice!

Assembly example: measuring cycles

- Time Stamp Counter
 - Special 64-bit register in Intel-compatible machines
 - Incremented every clock cycle
 - Read with `rdtsc` instruction
- Cannot be read from a programming language
 - Requires assembly code to access
 - C compiler's `asm` facility inserts assembly code into generated machine code

```
uint64_t rdtsc()
{
    uint32_t lo, hi;

    asm volatile("rdtsc; movl %%edx,%0; movl %%eax,%1"
        : "=r" (hi), "=r" (lo)
        :
        : "%edx", "%eax");
    return (lo | (((uint64_t)hi) << 32));
}
```

Assembly example: measuring cycles

- Can be used to measure *how many instruction cycles* a computation really took.
- Need to be careful about measurement methodology...

```
int main(int argc, char *argv[])
{
    uint64_t start, overhead;
    unsigned long result;

    // Measure the overhead.
    // We should really repeat this many times.
    start = rdtsc();
    overhead = rdtsc()-start;
    printf("Counter overhead is %lu cycles\n", overhead);

    // Time the function
    start = rdtsc();
    result = calc();
    printf("Time = %lu cycles\n", rdtsc()-start-overhead );
    printf("Result = %lu\n", result );
    return 0;
}
```


Inconvenient truth:

Memory is not a nice array that
stores your data

The details of memory

- Memory is **not unbounded**
 - It must be allocated and managed
 - Many applications are memory-dominated
- Memory performance is **not uniform**
 - Cache and virtual memory effects can greatly affect program performance
 - Adapting program to characteristics of memory system can lead to major speed improvements
- Memory is **typed**
 - Different kinds of memory behave differently

Memory-related bugs are still a nightmare.

```
typedef struct {
    int a[2];
    double d;
} struct_t;

double fun(int i) {
    volatile struct_t s;
    s.d = 3.14;
    s.a[i] = 1073741824; /* Possibly out of bounds */
    return s.d;
}
```

fun(0)	->	3.14
fun(1)	->	3.14
fun(2)	->	3.1399998664856
fun(3)	->	2.00000061035156
fun(4)	->	3.14
fun(6)	->	Segmentation fault

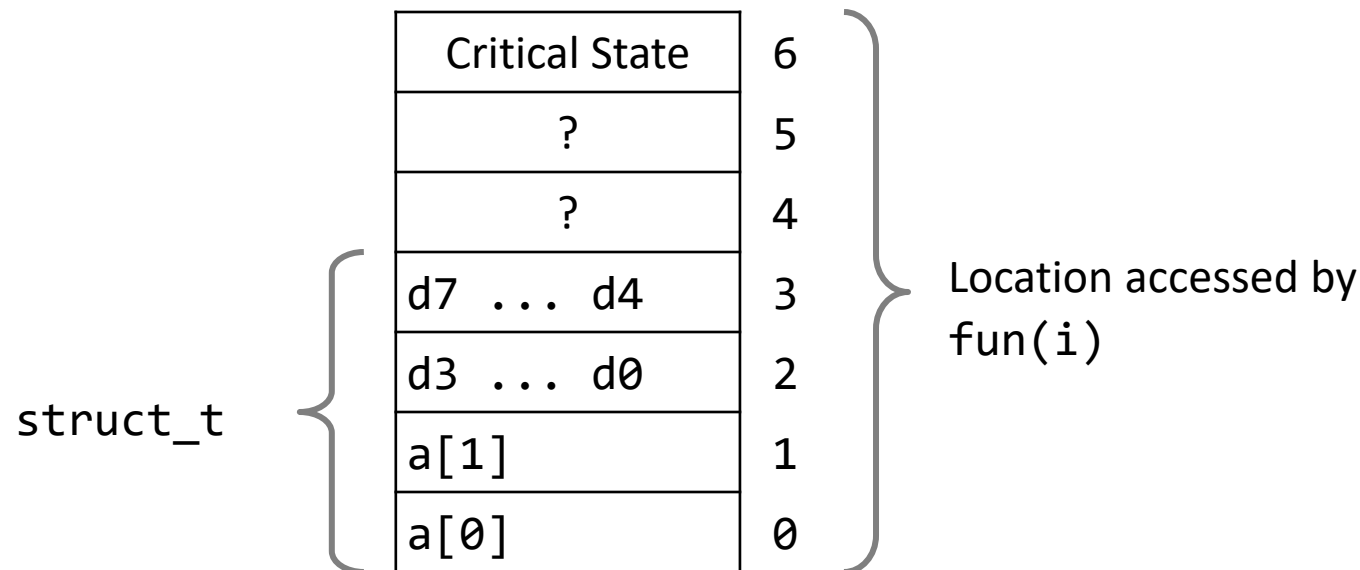
Actual results are
system-specific...

Memory referencing bug

```
typedef struct {  
    int a[2];  
    double d;  
} struct_t;
```

fun(0)	->	3.14
fun(1)	->	3.14
fun(2)	->	3.1399998664856
fun(3)	->	2.00000061035156
fun(4)	->	3.14
fun(6)	->	Segmentation fault

Explanation:



Memory system performance

```
void copyij(int src[2048][2048],
            int dst[2048][2048])
{
    int i,j;
    for (i = 0; i < 2048; i++)
        for (j = 0; j < 2048; j++)
            dst[i][j] = src[i][j];
}
```

5.2 ms


```
void copyji(int src[2048][2048],
            int dst[2048][2048])
{
    int i,j;
    for (j = 0; j < 2048; j++)
        for (i = 0; i < 2048; i++)
            dst[i][j] = src[i][j];
}
```

162 ms !

Intel Core i7 2.7 GHz

- Hierarchical memory organization
- Performance depends on access patterns
 - Including how step through multi-dimensional array

The Memory Mountain



Inconvenient truth:

Performance is about much more than
asymptotic complexity

There's much more to performance than asymptotic complexity

- **Constant factors** matter too – often more.
- Even **exact op count** does not predict performance
 - Easily see 10:1 performance range depending on how code written
 - Must optimize at multiple levels: algorithm, data representations, procedures, and loops
- Must **understand system** to optimize performance
 - How programs compiled and executed
 - How to measure program performance and identify bottlenecks
 - How to improve performance without destroying code modularity and generality

Example: matrix-matrix multiplication

- Fundamental operation in ML, graphics, etc., etc.

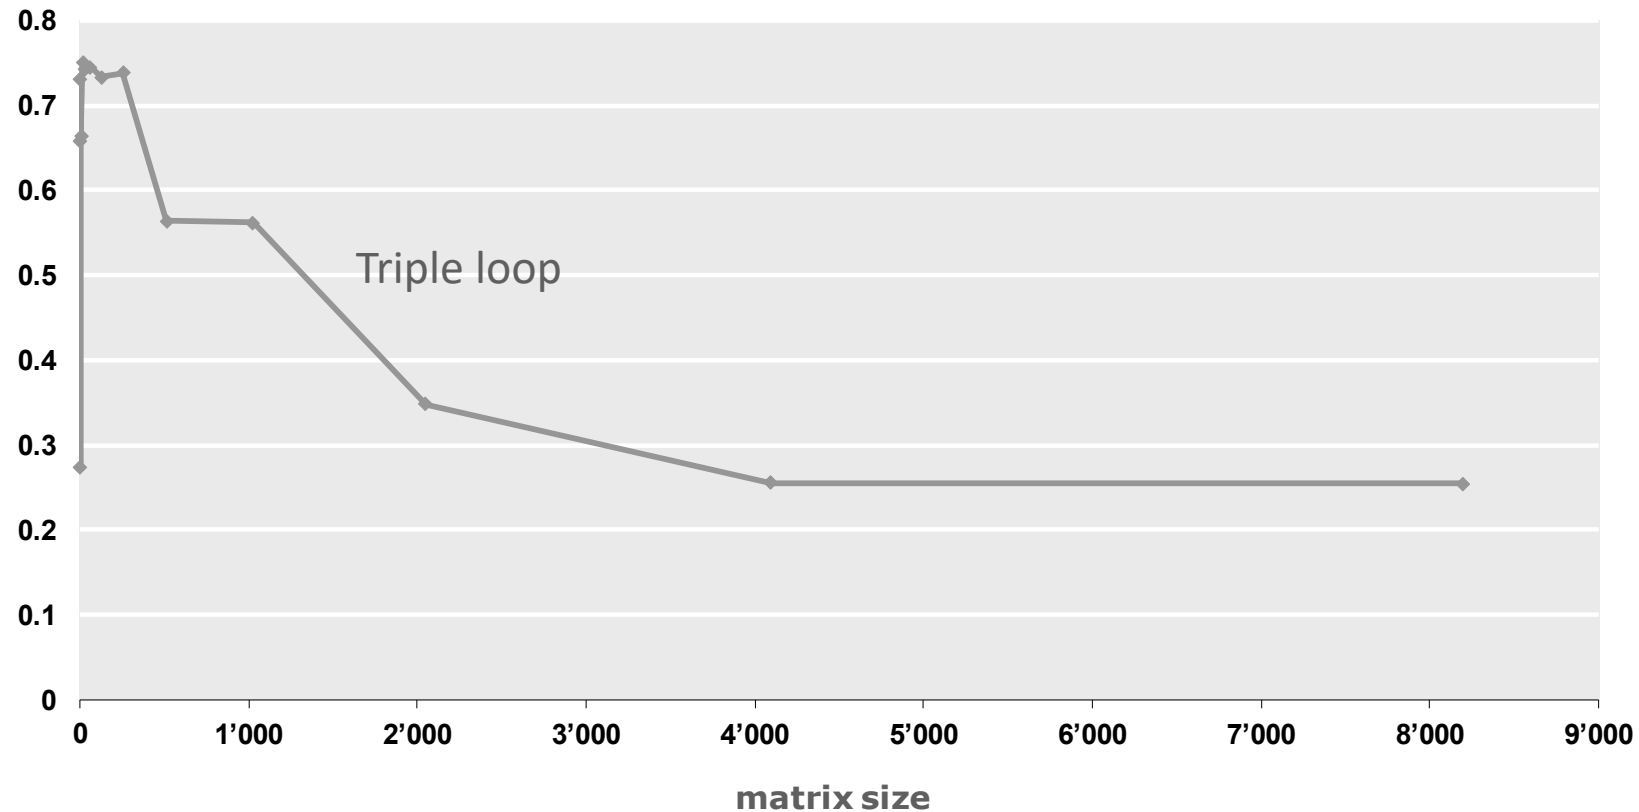
$$\left(\begin{smallmatrix} \dots \end{smallmatrix} \right) \leftarrow \left(\begin{smallmatrix} \dots \end{smallmatrix} \right) \times \left(\begin{smallmatrix} \dots \end{smallmatrix} \right)$$

- How complicated can this be?
- Basically requires n^3 operations for $n \times n$ matrices
 - *Though some subcubic algorithms exist – ignore for now.*


Example: matrix-matrix multiplication

Matrix-Matrix Multiplication (MMM) on 2 x Core 2 Duo 3 GHz (double precision)

Gflop/s




Example: matrix-matrix multiplication



- Standard desktop computer, vendor compiler, using optimization flags
- Both implementations have **exactly** the same operations count ($2n^3$)
- What is going on?

MMM plot: analysis

Matrix-Matrix Multiplication (MMM) on 2 x Core 2 Duo 3 GHz
Gflop/s



- Why? Blocking or tiling, loop unrolling, array scalarization, instruction scheduling, ...
- **Effect: less register spills, less L1/L2 cache misses, less TLB misses**

Inconvenient truth:

Computers don't just execute programs
Programs don't just calculate values

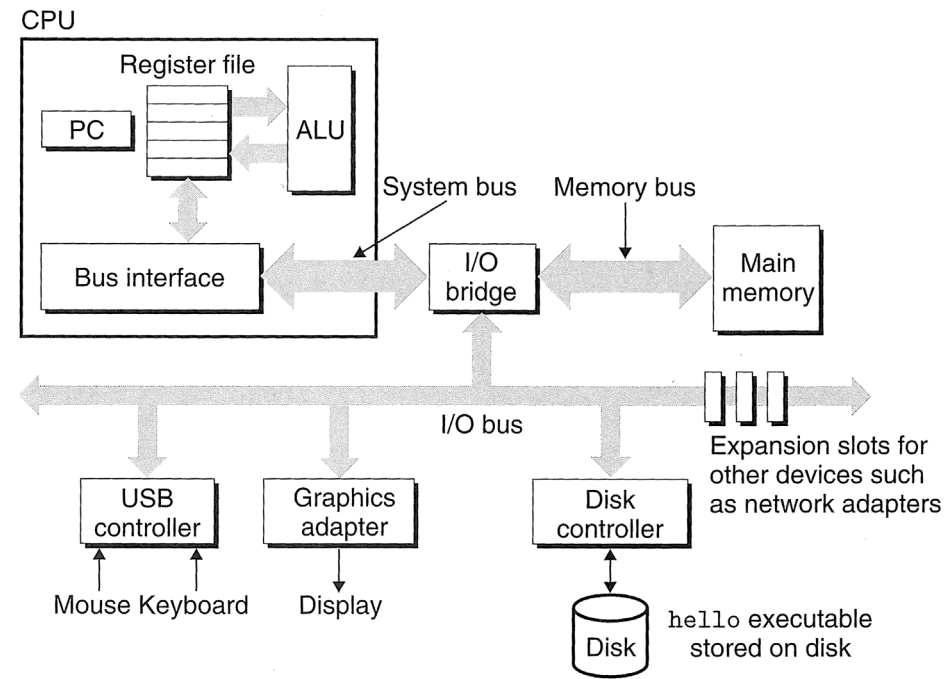
Computers don't just run programs

- They need to get data **in** and **out**
 - I/O critical to program reliability and performance
 - **Sense** the physical world
 - **Act** in the physical world
- They **communicate** over networks
 - Many system-level issues arise with a network
 - Concurrent operations by autonomous processes
 - Coping with unreliable media
 - Cross-platform interoperability
 - Complex performance issues

Lies our teachers tell us...

Figure 1.4

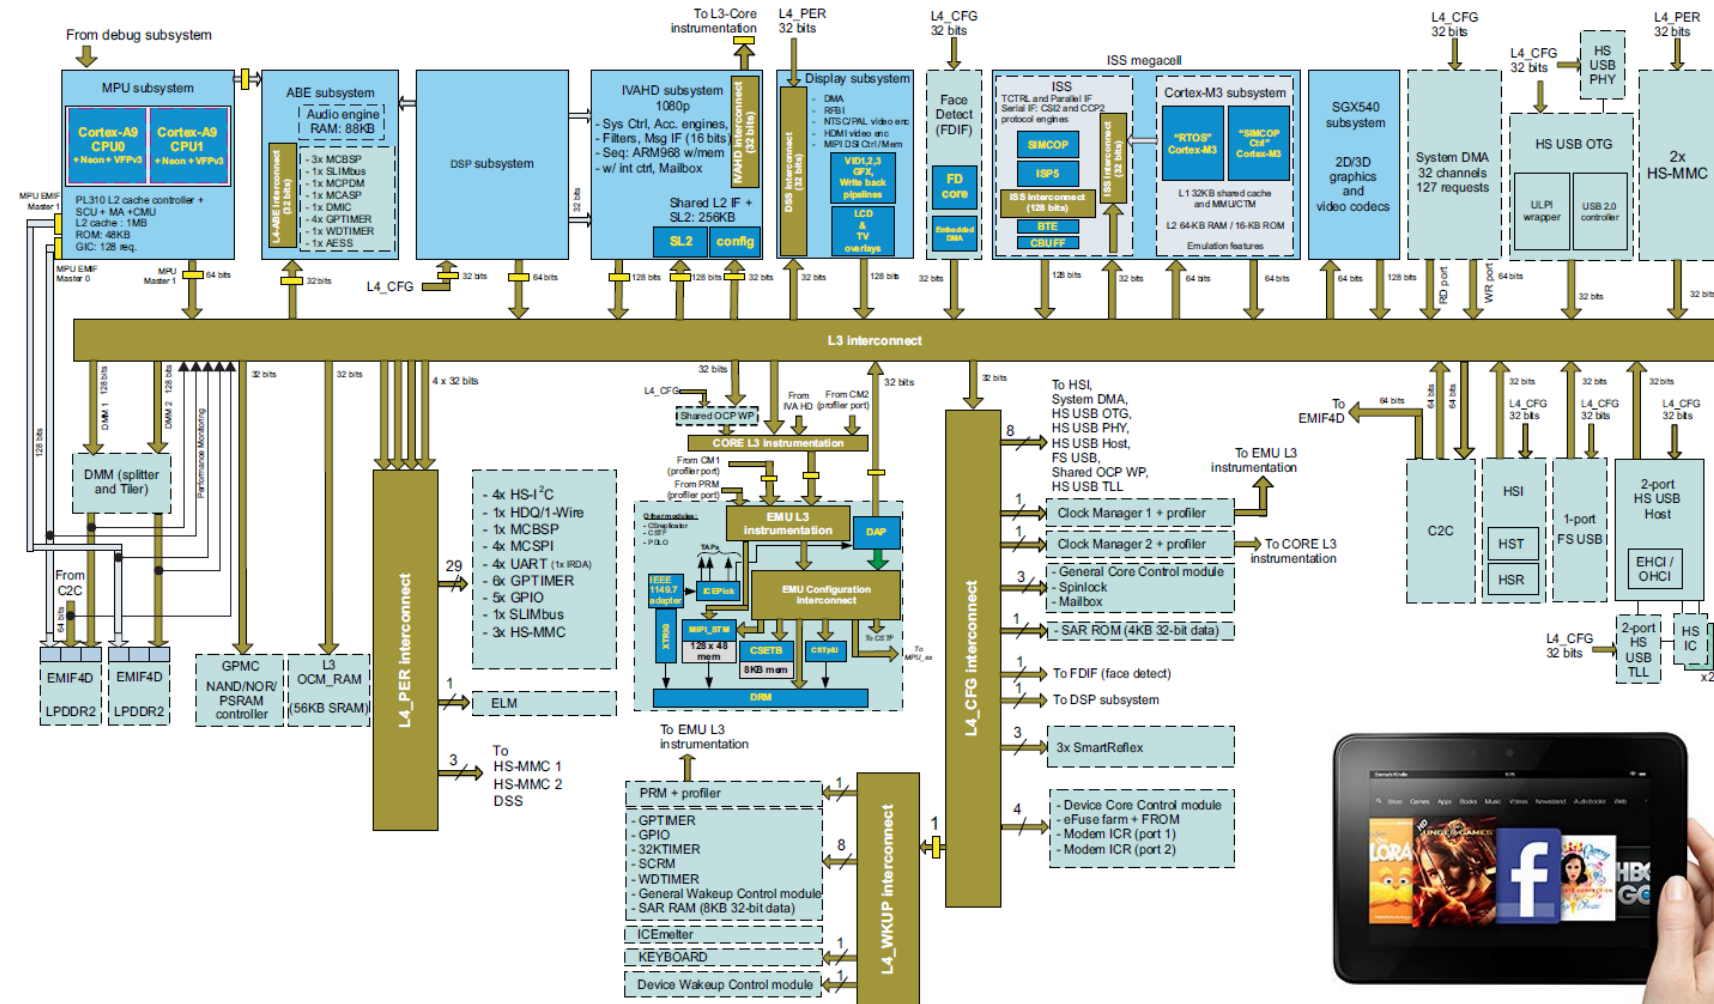
Hardware organization of a typical system. CPU: Central Processing Unit, ALU: Arithmetic/Logic Unit, PC: Program counter, USB: Universal Serial Bus.



systems, but all systems have a similar look and feel. Don't worry about the complexity of this figure just now. We will get to its various details in stages throughout the course of the book.

Computer Systems, A Programmer's Perspective, Bryant & O'Hallaron, 2011

A modern(ish) System-on-Chip



Texas Instruments
OMAP 4460, c.2011



Inconvenient truth:

Programs are not semantic
specifications

The role of “standards”

- Language standards aim to **specify unambiguously** what any program in the language does when compiled and executed.
 - Java: “write once, run anywhere”
 - Formal semantics
- The C standards should be viewed as rather different
 - Behavior frequently described as “*implementation dependent*”
 - What does this mean?

“Implementation defined”

“unspecified behavior where each implementation documents how the choice is made”

At least two options:

- Compiler is allowed to do **anything**, so **optimizes out** the code completely
- Compiler implements the **most natural mapping** to the target hardware and **documents** this.

3. Terms, definitions, and symbols

For the purposes of this International Standard, the following definitions apply. Other terms are defined where they appear in *italic* type or on the left side of a syntax rule. Terms explicitly defined in this International Standard are not to be presumed to refer implicitly to similar terms defined elsewhere. Terms not defined in this International Standard are to be interpreted according to ISO/IEC 2382-1. Mathematical symbols not defined in this International Standard are to be interpreted according to ISO 31-11.

3.1

access

⟨execution-time action⟩ to read or modify the value of an object

NOTE 1 Where only one of these two actions is meant, “read” or “modify” is used.

NOTE 2 “Modify” includes the case where the new value being stored is the same as the previous value.

NOTE 3 Expressions that are not evaluated do not access objects.

3.2

alignment

requirement that objects of a particular type be located on storage boundaries with addresses that are particular multiples of a byte address

3.3

argument

actual argument

actual parameter (deprecated)

expression in the comma-separated list bounded by the parentheses in a function call expression, or a sequence of preprocessing tokens in the comma-separated list bounded by the parentheses in a function-like macro invocation

3.4

behavior

external appearance or action

3.4.1

implementation-defined behavior

unspecified behavior where each implementation documents how the choice is made

EXAMPLE An example of implementation-defined behavior is the propagation of the high-order bit when a signed integer is shifted right.

3.4.2

locale-specific behavior

behavior that depends on local conventions of nationality, culture, and language that each implementation documents

“Implementation defined”

“unspecified behavior where each implementation documents how the choice is made”

At least two interpretations:

- Compiler is allowed to do **anything**, so **optimizes out** the code completely
- Compiler implements the **most natural mapping** to the target hardware and **documents** this.

3. Terms, definitions, and symbols

For the purposes of this International Standard, the following definitions apply. Other terms are defined where they appear in *italic* type or on the left side of a syntax rule. Terms explicitly defined in this International Standard are not to be presumed to refer implicitly to similar terms defined elsewhere. Terms not defined in this International Standard are to be interpreted according to ISO/IEC 2382-1. Mathematical symbols not defined in this International Standard are to be interpreted according to ISO 31-11.

3.1

Default behavior for newer C compilers ☹️

argument

actual argument

actual parameter (deprecated)

expression in the comma-separated list bounded by the parentheses in a function call expression, or a sequence of preprocessing tokens in the comma-separated list bounded by the parentheses in a function-like macro invocation

3.4

behavior

external appearance or action

3.4.1

implementation-defined behavior

unspecified behavior where each implementation documents how the choice is made

EXAMPLE An example of implementation-defined behavior is the propagation of the high-order bit when a signed integer is shifted right.

3.4.2

locale-specific behavior

behavior that depends on local conventions of nationality, culture, and language that each implementation documents

“Implementation defined”

“unspecified behavior where each implementation documents how the choice is made”

At least two interpretations:

- Compiler is allowed to do **anything**, so **optimizes out** the code completely
- Compiler implements the **most natural mapping** to the target hardware and **documents** this.

3. Terms, definitions, and symbols

For the purposes of this International Standard, the following definitions apply. Other terms are defined where they appear in *italic* type or on the left side of a syntax rule. Terms explicitly defined in this International Standard are not to be presumed to refer implicitly to similar terms defined elsewhere. Terms not defined in this International Standard are to be interpreted according to ISO/IEC 2382-1. Mathematical symbols not defined in this International Standard are to be interpreted according to ISO 31-11.

3.1

Default behavior for newer C compilers ☹️

Default behavior for older C compilers, and *What You Actually Want.*

The role of “standards”

- Language standards aim to **specify unambiguously** what any program in the language does when compiled and executed.
 - Java: “write once, run anywhere”
 - Formal semantics
- The C standards should be viewed as rather different
 - Behavior frequently described as “**implementation dependent**”
 - What does this mean?


A program is a set of instructions to a compiler that tell it what assembly language to generate.

Summary

1. `ints` are not integers. `floats` are not real numbers.
2. You've got to know assembly.
3. Memory matters. RAM is an unrealistic abstraction.
4. There's much more to performance than asymptotic complexity.
5. Computers don't just evaluate programs.
6. Programs are not complete semantic specifications.


Interesting times...

- Processors are not getting faster.
- Performance-wise, progress in computer architecture has halted.
 - In fact, it's going backwards due to security concerns.




Interesting times...

- Lots of companies, universities, etc. are building new kinds of computers.
- How can these be programmed?
- This is a systems software problem.



Computers are looking different!

- Cavium / Marvell ThunderX-1 NP:
 - 48 x ARMv8 cores at 2GHz
 - 128 GB DDR4
 - 2 x 40Gb/s network
- Xilinx UltraScale+ VU9P
 - 512 GB DDR4
 - 4 x 100 Gb/s network
- NVMe, SATA, PCIe on both sides
- Native coherence between FPGA and CPU



1.4: What we'll assume you know

Systems Programming and Computer Architecture

Courses already

- Programming & software engineering
- Parallel programming
- Data structures and algorithms
- Digital Circuits
- Discrete Mathematics

What we'll assume you know #1:


- Memory, addresses, bytes, and **words**
 - Binary, and Hexadecimal notation
 - Byte-ordering (Big/Little Endian)
- Boolean algebra
 - and, or, not, xor
 - Generalized Boolean algebra: bitwise operations on words as bit vectors
- How to write programs
- Languages
 - Java (but we won't use it)
 - Some C or C++ (but not much)
 - Some assembly language

What we'll assume you know #2:

- Processor architecture, pipelines
 - Registers
 - Addressing modes
 - Instruction formats
- Software engineering
 - Object-orientation
 - Design-by-contract
 - Strong typing
- Basic memory systems
 - cache architectures
 - virtual memory
 - I/O devices
- Concurrency and parallelism
 - Threads
 - Locks, mutexes, condition variables
 - Parallel programming constructs



Textbooks

- Randal E. Bryant and David R. O'Hallaron,
 - Computer Systems: A Programmer's Perspective, Third Edition (CS:APP3e), Pearson, 2016
 - <http://csapp.cs.cmu.edu>




Books on C (there are many)

- Brian Kernighan and Dennis Ritchie,
 - “The C Programming Language, Second Edition”, Prentice Hall, 1988(!)



- Samuel Harbison and Guy Steele
 - C: A Reference Manual
 - 5th edition 2002

- Peter van der Linden
 - Expert C Programming: Deep C Secrets, 1994



OK, let's start