



Chapter 2: Introduction to C

252-0061-00 V Systems Programming and Computer Architecture

Goal

- Introduction to C
 - Enough to program assignments
 - Background for lectures
- Assume you know Java (or C#)
 - E.g. from Parallel Programming
- Non-goal (for now...):
 - Teach details and strict definition of C
 - Teach advanced features/idioms/techniques in C
 - But still much more than you saw last year...

2.1: History and toolchain

Systems Programming and Computer Architecture

History

- Developed 1969-1972 by Dennis Ritchie
 - CPL → BCPL → B → C
 - Highly influenced by DEC PDP-11 architecture
 - Portable across many architectures
- Standards:
 - K&R C (standard was the compiler source!)
 - ANSI C
 - **C99 (we'll use this)**
 - C11 (more recent, less used)
 - C17 (even more recent, bug fixes to C11...)
 - ... and many C-like variants



1941-2011

Compared to Java, C#, PHP, Python, etc...

- Very **fast**
 - Almost impossible to write assembly as fast as a good C compiler
 - Pretty much impossible to compile Java to run as fast as C
- Powerful macro **pre-processor** (cpp)
- Close to the metal: you can *know* what the code is doing to the hardware

⇒ Remains the language of choice for

- Operating System developers
- Embedded systems
- People who *really* care about speed
- Authors of security exploits

Just some of what you don't get

- No **objects, classes, traits, features, methods, or interfaces**
 - Only functions/procedures
 - We will see function pointers later...
- No fancy built-in **types**
 - Mostly just what the hardware provides
 - Type constructors to build structured types
- No **exceptions**
 - convention is to use integer *return codes*

Most important difference

- No automatic **memory management**
 - Lots of things on the stack
 - No garbage collection
 - Heap structures explicitly created and freed
- **Pointers**: direct access to memory addresses
 - Weakly typed by what they point to

C is about **directly** building and manipulating structures in **main memory**!

Syntax: the good news

- Java, JavaScript, C++, and C# **syntax** almost entirely lifted from C
- **Comments** (`/*...*/`, `//`) the same
- **Identifiers** same as in Java
 - C# allows more characters in identifiers
- **Block** structure using `{ ... }`

Many other constructs the same or similar.

Syntax: main differences

- List of **reserved words** is different
- C is run through a ***macro preprocessor***
 - String and file substitution
 - Conditional compilation
 - Although C# has preprocessor directives, it does not have a separate preprocessor.
Moreover there are no macros.

Hello world

```
#include <stdio.h>

int main(int argc, char *argv[])
{
    printf("hello, world\n");
    return 0;
}
```

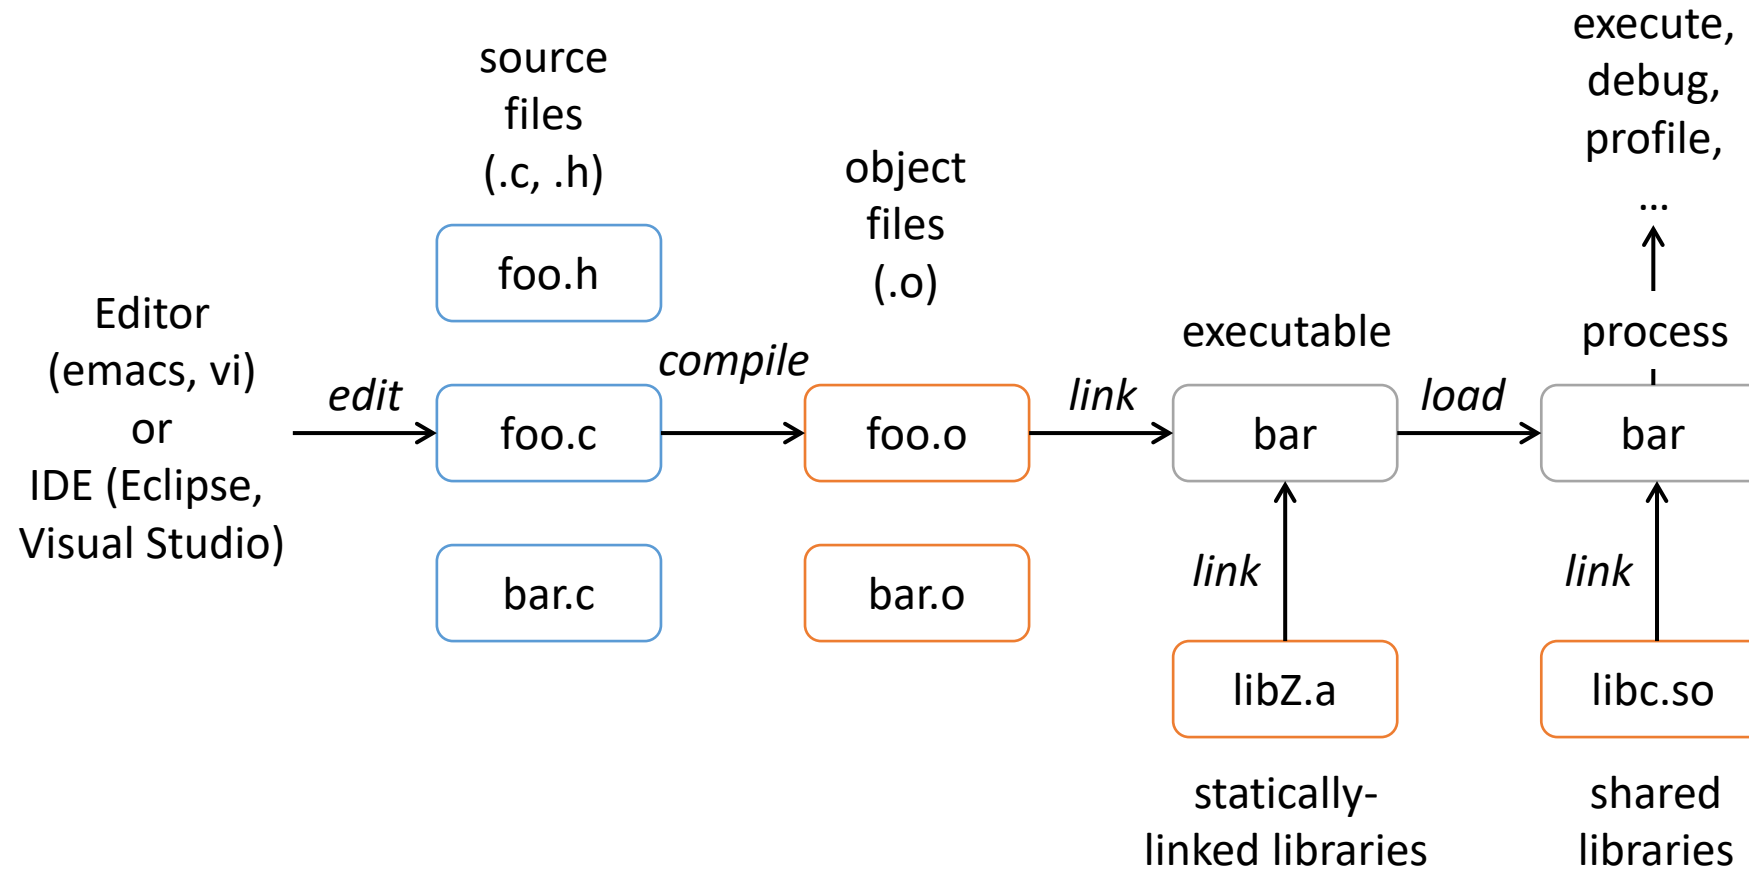
“header file” – bit like an interface file in Java or C#

Every program has to have a “main” function, which takes a list of *command line arguments*.

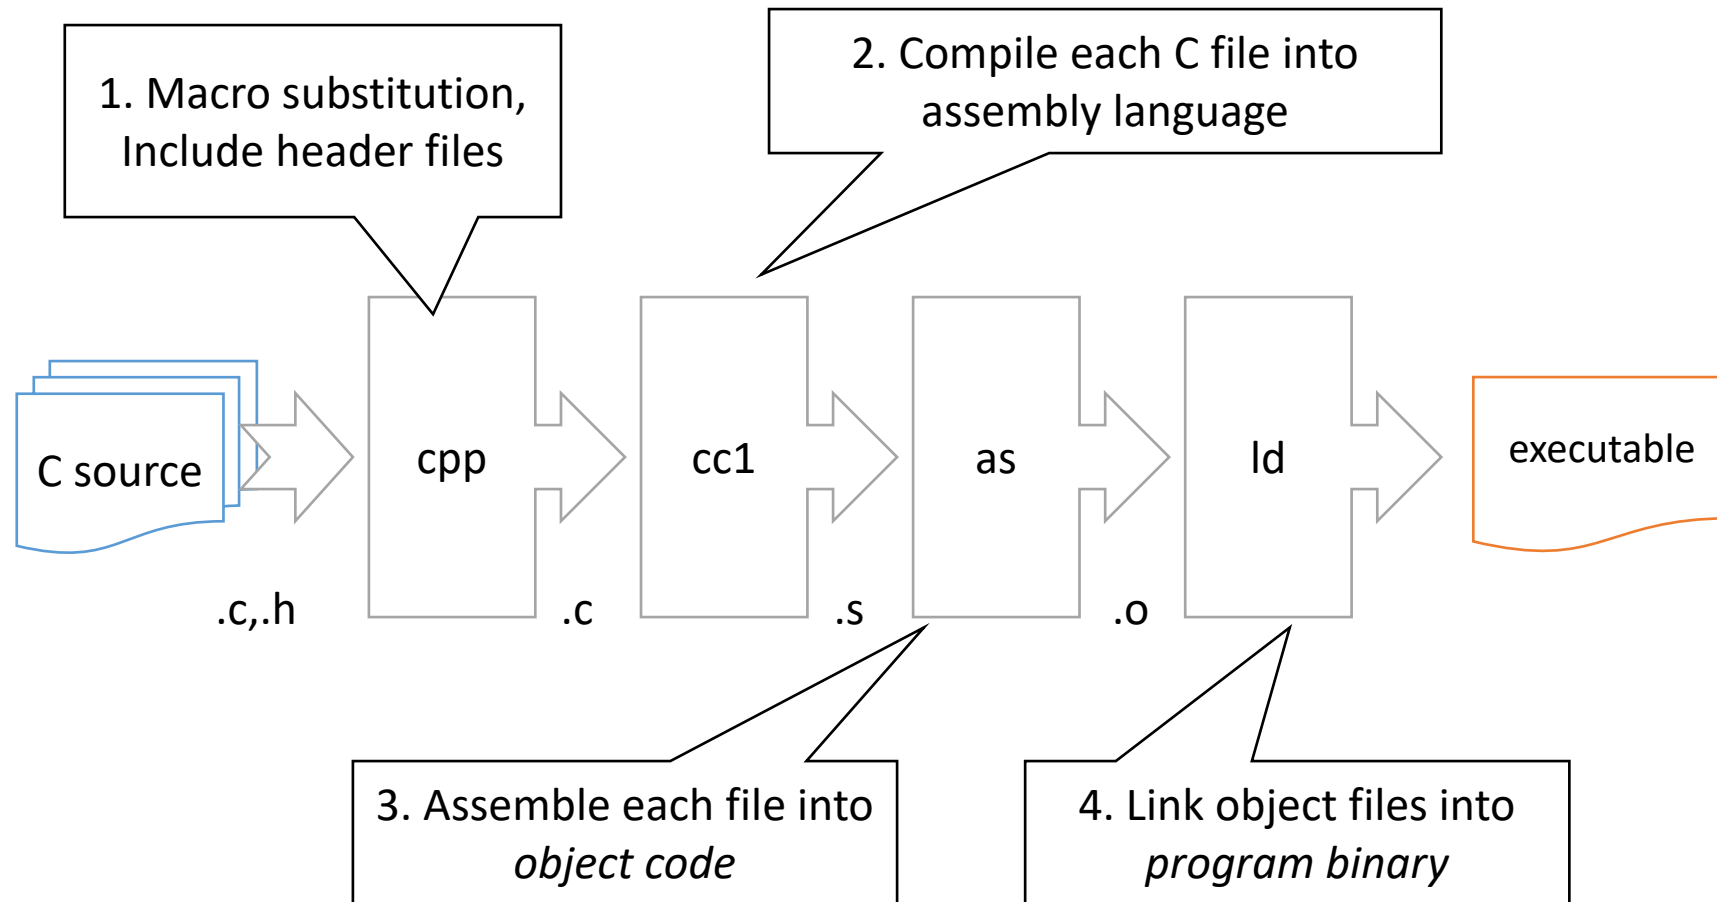
Generic function for printing formatted strings. The “newline” is not included automatically!

Returning 0 indicates everything is OK – C has no exceptions.

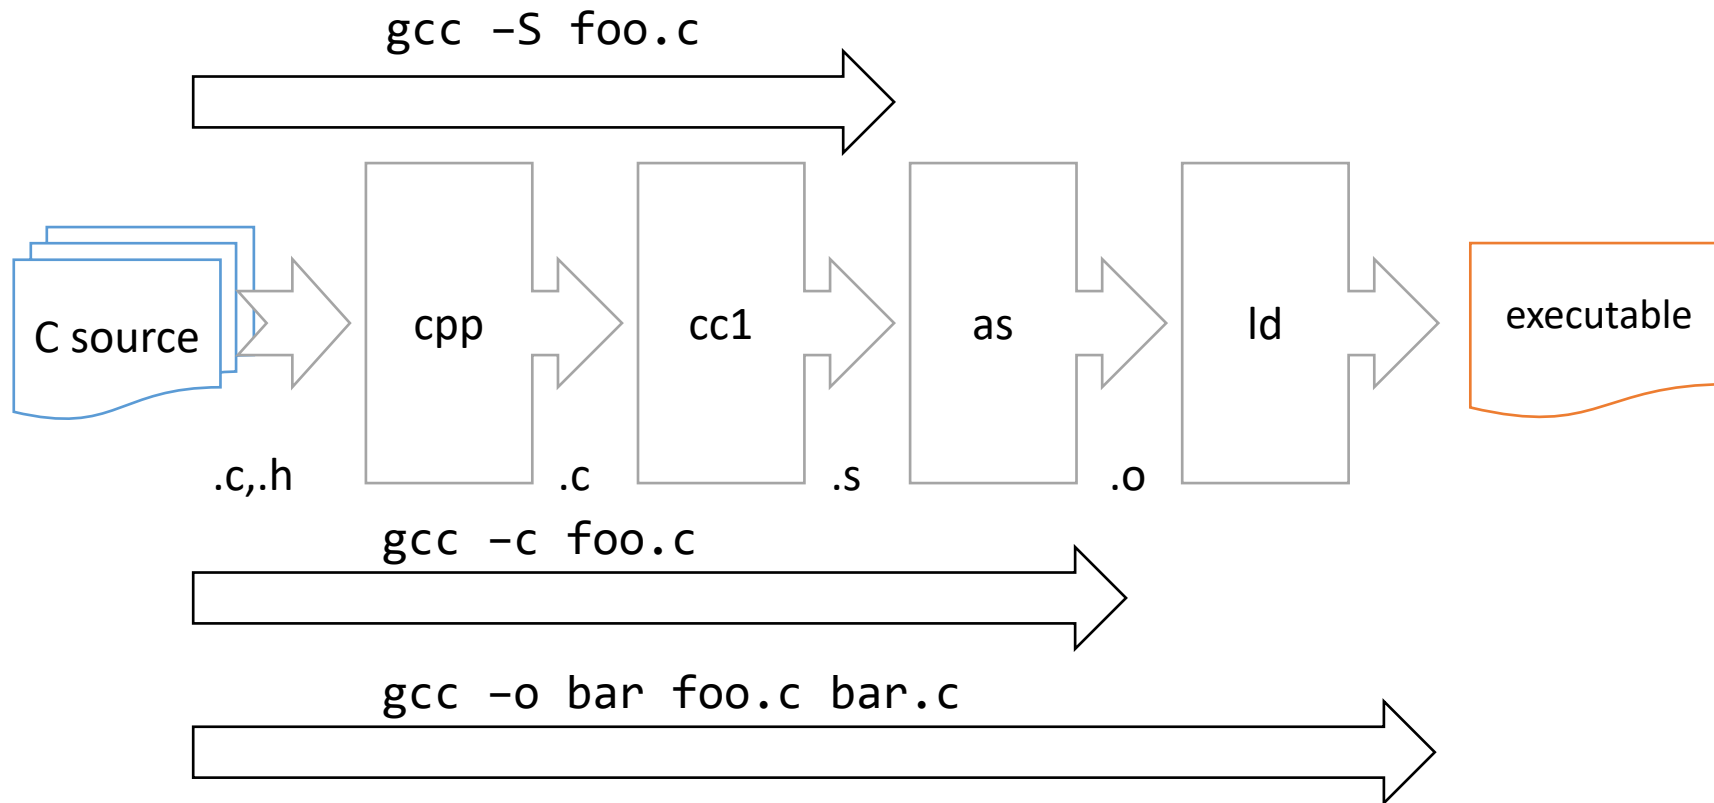
Workflow



GNU gcc Toolchain



GNU `gcc` Toolchain



Summary

- C is a *systems programming language!*
 - It is there to program the system.
 - Also useful for high performance
 - Understanding C is about understanding how
 - Your program
 - The C compiler
 - The present computer system
- ... all interact with each other.

2.2: Control flow in C

Systems Programming and Computer Architecture

Control flow statements (like Java or C# or C++)

```
if (Expression) Statement_when_true  
    else Statement_when_false
```

```
switch (Expression) {  
    case Constant_1 : Statement; break;  
    case Constant_2 : Statement; break;  
    ...  
    case Constant_n : Statement; break;  
    default: Statement; break;  
}
```

```
return (Expression)
```

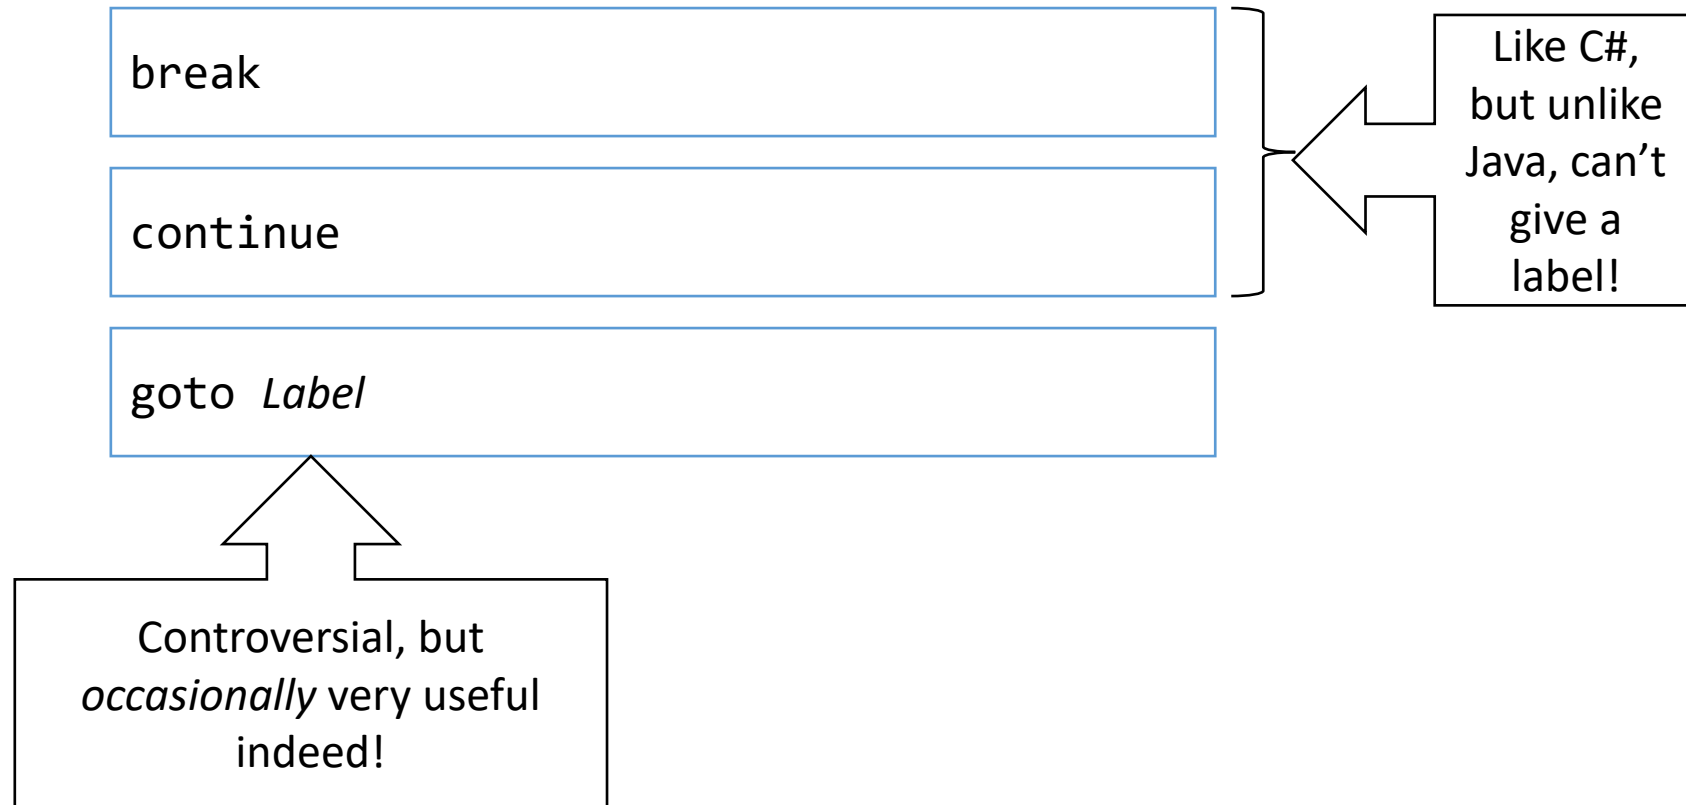

Control flow statements (like Java or C# or C++)

```
for (Initial; Test; Increment) Statement
```

```
while (Expression) Statement
```

```
do Statement while (Expression)
```

Control flow statements (not like Java, same as C#)



Functions: similar to Java

- Name
- Return type
- Argument types
- Body

```
#include <stdio.h>

// Compute factorial function
// fact(n) = n * (n-1) * ... * 2 * 1
int fact(int n)
{
    if (n == 0) {
        return(1);
    } else {
        return(n * fact(n-1));
    }
}

int main(int argc, char *argv[])
{
    int n, m;

    printf("Enter a number: ");
    scanf("%d", &n);
    m = fact(n);
    printf("Factorial of %d is %d.\n", n, m);
    return 0;
}
```

main(): also a function

```
#include <stdio.h>

int main(int argc, char *argv[])
{
    int i;

    // Print arguments from command line */
    printf("argc = %d\n\n", argc);
    for (i = 0; i < argc; ++i) {
        printf("argv[%d]: %s\n", i, argv[i]);
    }
    return 0;
}
```

Basic I/O: printf()

- Just another function, but very useful!

```
#include <stdio.h>

int main(int argc, char *argv[])
{
    int i = 314;
    const char s[] = "Mothy";
    printf("My name is %s and I work in STF H %d\n", s, i);
    return 0;
}
```

printf is
a *variadic*
function!

- First argument is format string
 - see “man 3 printf” for all the (many) options
- Remaining arguments are arbitrary
 - but must match the format
- You will see other “printf-like” functions

Summary: control flow in C

- Functions
 - `return (..)`
- Loops
 - `for(..; ..; ..)`
 - `do .. while (..)`
 - `while (..) ..`
- Conditionals
 - `if (..) then .. else ..`
 - `switch (..) case .. : ..; default ..`
- Jumps
 - `break, continue`
 - `goto ..`
- I/O:
 - `printf()`

2.3: Basic types in C

Systems Programming and Computer Architecture

Declarations

- Are like Java:
 - `int my_int;`
 - `double some_fp = 0.123;`
- Inside a block:
 - Scope is just the block
 - `static` → value *persists* between calls
- Outside a block:
 - Scope is the *entire program!*
 - `static` → scope limited to the file (*compilation unit*)

```
#include <stdio.h>

static int j = 0;

int func(int j)
{
    static int i = 0;

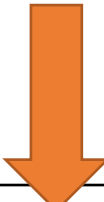
    i = i + 1;
    j = j + 1;
    printf("In func: i=%d, j=%d\n", i, j);
    return j;
}

int main(int argc, char *argv[])
{
    int i = 0;

    printf("In main: i=%d, j=%d\n", i, j);
    func(j);
    printf("In main: i=%d, j=%d\n", i, j);
    func(j);
    printf("In main: i=%d, j=%d\n", i, j);
    return 0;
}
```


Integers and floats

- Types and sizes:



| C data type | Typical 32-bit | ia32 | Intel x86-64 |
|-------------|----------------|-------|--------------|
| char | 1 | 1 | 1 |
| short | 2 | 2 | 2 |
| int | 4 | 4 | 4 |
| long | 4 | 4 | 8 |
| long long | 8 | 8 | 8 |
| float | 4 | 4 | 4 |
| double | 8 | 8 | 8 |
| long double | 8 | 10/12 | 10/16 |

Sizes are
implementation
defined!

- Integers are **signed** by default
 - use signed or unsigned to clarify

C99 extended integer types

```
#include <stdint.h>
```

```
int8_t      a;
```

```
int16_t     b;
```

```
int32_t     c;
```

```
int64_t     d;
```

Signed integers,
precise size in bits

```
uint8_t     x;
```

```
uint16_t    y;
```

```
uint32_t    z;
```

```
uint64_t    w;
```

Unsigned integers,

Integers and floats

- Rules for arithmetic on integers and floats are **complicated!**
 - Implicit conversions between integer types
 - Implicit conversions between floating point types
 - Explicit conversions between anything (casts)
- Behavior is either:
 - Determined by the hardware (*implementation defined*)
 - Was decided by hardware, a long time ago (standardized)
- We'll cover this more later...

Booleans

- Historically, boolean values are just integers
 - False → zero
 - True → anything non-zero
 - Negation (“!”) turns zero into non-zero, and vice-versa
- C99: new `bool` type supported
 - Completely optional, it’s still an integer
 - `#include <stdbool.h>`

Booleans

- Any statement in C is also an expression, hence:

```
#include <stdio.h>
#include <stdlib.h>

#define ERR_OUT_OF_RANGE (-1)

int test(int i)
{
    if (i >= 0 && i < 10) {
        printf("Success: value is %d\n", i);
        return 0;
    } else {
        return ERR_OUT_OF_RANGE;
    }
}

int main(int argc, char *argv[])
{
    int rc;

    if (argc != 2) {
        fprintf(stderr, "Usage: %s <number>\n", argv[0]);
        return 1;
    }
    if ( (rc = test(atoi(argv[1]))) ) {
        fprintf(stderr, "Error: argument was out of range\n");
        return 1;
    }
    return 0;
}
```



Booleans

- Any statement in C is also an expression, hence:

```
#include <stdio.h>

#define ERR_OUT_OF_RANGE (-1)

int main(int argc, char *argv[])
{
    FILE *f;

    if (argc != 2) {
        fprintf(stderr, "Usage: %s <filename>\n", argv[0]);
        return 1;
    }
    if (!(f = fopen( argv[1], "r" ))) {
        perror("Failed to open file");
        return 1;
    }
    printf("Successfully opened %s\n", argv[0]);
    fclose(f);
    return 0;
}
```

void

- There is a type called `void`.
- It has **no** value.
- Used for:
 - Untyped pointers (to raw memory):
 `"void *"`
 - Declaring functions with no return value (procedures)

Summary: C basic types

- Declarations
- Scopes and `static`
- Integers and floats, extended types
- Booleans
- `void`

2.4: Operators

Systems Programming and Computer Architecture

Operators: similar to Java

| Decreasing precedence ↓ | Operator | Associativity | |
|----------------------------|-----------------------------------|---------------|---|
| | () [] -> . | Left-to-right | <ul style="list-style-type: none">• () here is a function call• -> means <i>struct pointer indirection</i> |
| | ! ~ ++ -- + - * & (type) sizeof | Right-to-left | |
| | * / % | Left-to-right | |
| | + - | Left-to-right | <ul style="list-style-type: none">• Unary +, -, *• * here is <i>pointer indirection</i> |
| | << >> | Left-to-right | |
| | < <= > >= | Left-to-right | |
| | == != | Left-to-right | |
| | & | Left-to-right | |
| | ^ | Left-to-right | |
| | | Left-to-right | |
| | && | Left-to-right | <ul style="list-style-type: none">• Ternary if-else operator |
| | | Left-to-right | |
| | ?: | Right-to-left | <ul style="list-style-type: none">• Assignment operators |
| | = += -= *= /= %= &= ^= = <<= >>= | Right-to-left | |
| | , | Left-to-right | |

Assignment operators

- In many imperative languages

`x = foo();`

- is an assignment **statement**.

- In C, it is an **expression**!
 - Value is the value being assigned

- Also:

`x += y` \leftrightarrow `x = x + y`

- and so with `-=`, `*=`, `<<=`, etc.

What is associativity again?

- Left-to-right associativity:

- $A + B + C \rightarrow (A + B) + C$
- $A + B + C + D \rightarrow ((A + B) + C) + D$

- Right-to-left:

- $A += B += C \rightarrow A += (B += C)$
- Makes sense here, but elsewhere it's rarely what you want...

Post-increment and pre-increment

- `i++`
 - Value: current value of `i`
 - Effect: $i \leftarrow i+1$
- `++i`
 - Effect: $i \leftarrow i+1$
 - Value: new value of `i`
- Conversely `i--` and `--i`
- Works for any scalar type
 - Importantly: works for pointers!

Historical:

Digital PDP computers had pre- and post-increment and -decrement addressing modes



Casting

- Most C types can be *cast* to another:

```
unsigned int ui = 0xDEADBEEF;  
signed int i    = (signed int)ui;
```

⇒ i has value -559038737.

Name of type
in parentheses
functions like an
operator.

- Bit-representation does (usually...) not change
- Frequently used with pointer types...

Summary: C operators

- Operators and precedence
- Assignment operators
- Post/pre inc/decrement
- Casting

2.5: Arrays in C

Systems Programming and Computer Architecture

Arrays

- Finite vector of variables, all the same type
- For an N-element array a:
 - First element is a[0]
 - last element is a[N-1]
- C compiler **does not** check the array bounds!
 - Very typical bug!
 - Always check array bounds!

```
#include <stdio.h>

float data[5]; // data to average and total
float total;   // total of the data items
float average; // average of the items

int main() {
    data[0] = 34.0;
    data[1] = 27.0;
    data[2] = 45.0;
    data[3] = 82.0;
    data[4] = 22.0;

    total = data[0] + data[1] + data[2] + data[3] + data[4];
    average = total / 5.0;
    printf("Total %f Average %f\n", total, average);
    return (0);
}
```


Multi-dimensional arrays

```
int a[3][3];
```

| | | | | | | | | |
|---------|-----|---------|-----|---------|-----|---------|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| a[0][0] | ... | a[1][0] | ... | a[1][2] | ... | a[2][2] | | |

```
int a = 1;
for (i=0; i < 3; i++)
    for (j=0; j < 3; j++)
        matrix[i][j] = a++;
```

| | | | | | | | | |
|---|---|---|---|---|---|---|---|---|
| 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 |
|---|---|---|---|---|---|---|---|---|



```
int a = 1;
for (i=0; i < 3; i++)
    for (j=0; j < 3; j++)
        matrix[j][i] = a++;
```

| | | | | | | | | |
|---|---|---|---|---|---|---|---|---|
| 1 | 4 | 7 | 2 | 5 | 8 | 3 | 6 | 9 |
|---|---|---|---|---|---|---|---|---|



Array initializers

- Arrays can be initialized when they are defined:

```
#include <stdio.h>

int main(int argc, char *argv[])
{
    int i, j;
    int a[3] = {3, 7, 9};
    int b[3][3] = {
        { 1, 2, 3 },
        { 4, 5, 6 },
        { 7, 8, 9 },
    };

    for(i = 0; i < 3; i++) {
        printf("a[%d] = %d\n", i, a[i] );
        for(j = 0; j < 3; j++) {
            printf("  b[%d][%d] = %d\n", i, j, b[i][j] );
        }
    }

    return 0;
}
```

Array initializers

- Arrays can be initialized when they are defined:

```
#include <stdio.h>

#define ARRAY_SIZE 10

int main(int argc, char *argv[])
{
    int i;

    float lu[ARRAY_SIZE];      // Uninitialized
    float li[ARRAY_SIZE] = {}; // Initialized to 0.0

    for(i = 0; i < ARRAY_SIZE; i++) {
        printf("li[%d] = %f, lu[%d] = %f\n", i, li[i], i, lu[i]);
    }
    return 0;
}
```

Strings

- C has no real string type!
 - Instead...
- Array of char's terminated with null byte
 - `0` or `'\0'`

```
#include <stdio.h>

int main(int argc, char *argv[])
{
    int i;

    // These strings are identical
    char s1[6] = "hello";
    char s2[6] = { 'h', 'e', 'l', 'l', 'o', 0 };

    printf("s1 = '%s'\n", s1);
    printf("s2 = '%s'\n", s2);

    for( i = 0; i < 6; i++ ) {
        printf("s1[%d] = %d, s2[%d] = %d\n", i, s1[i], i, s2[i]);
    }
    return 0;
}
```

String library functions

Generally named 'str**xxx**()'

```
#include <stdio.h>
#include <string.h>

int main(int argc, char *argv[])
{
    char name1[12], name2[12];
    char mixed[25], title[20];

    strcpy(name1, "Rosalinda");
    strcpy(name2, "Zeke");
    strcpy(title, "This is the title.");

    printf("      %s\n\n", title);
    printf("Name 1 is %s\n", name1);
    printf("Name 2 is %s\n", name2);

    ...
}
```

```
...

/* returns 1 if name1 > name2 */
if (strcmp(name1, name2) > 0) {
    strcpy(mixed, name1);
} else {
    strcpy(mixed, name2);
}
printf("The biggest name alphabetically is %s\n", mixed);

strcpy(mixed, name1);
strcat(mixed, " ");
strcat(mixed, name2);
printf("Both names are %s\n", mixed);
return 0;
}
```

This is the title.

Name1 is Rosalinda
 Name2 is Zeke
 The biggest name
 alphabetically is Zeke
 Both names are Rosalinda
 Zeke

Summary: C arrays

- Arrays of basic types
- Multidimensional arrays
- Initializers
- Strings
 - Arrays of ASCII characters
 - Null-terminated
- String library

We'll see **a lot** more
C as the course
progresses...